



Container Orchestration using Kubernetes Part - 2


Week 4

Topics

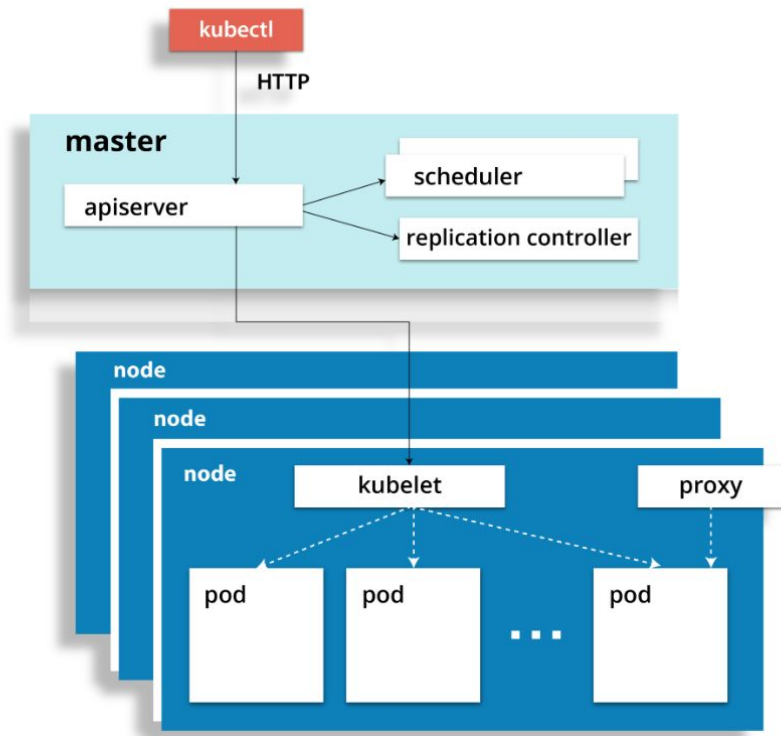
- **Part - 1**
- Container Orchestration Concept
- Introduction to Kubernetes
- Features and Advantages of Kubernetes
- Creating a Kubernetes Cluster
- **Part - 2**
- Deployment in Kubernetes
- Services in Kubernetes
- Rolling updates in Kubernetes
- Kubernetes Dashboard



Kubernetes Cluster Concept



Deploying an Application



- In order to deploy an application we can use “kubectl” command
- A proxy is an agent that helps pods to communicate with each other and other components inside a cluster.
- We need to install a network for proxy to work. We can use Calico or Flannel.
- Kubelet is an agent that runs on every node in a cluster
 - It is responsible for managing pods and their containers
 - Deals with pod specs YAML

Creating a New Pod - Imperative

- Creating Pods using imperative commands:

- kubectl run master-ac-demo --image=nginx

```
asif@master-node:~$ kubectl run master-ac-demo --image=nginx
pod/master-ac-demo created
asif@master-node:~$ kubectl get pods
NAME          READY   STATUS             RESTARTS   AGE
master-ac-demo 0/1     ContainerCreating   0           9s
asif@master-node:~$ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
master-ac-demo 1/1     Running   0           22s
asif@master-node:~$
```

- kubectl get pods -o wide

```
asif@master-node:~$ kubectl get pods -o wide
NAME          READY   STATUS    RESTARTS   AGE   IP          NODE      NOMINATED NODE   READINESS GATES
master-ac-demo 1/1     Running   0           11m   10.244.2.3   worker02   <none>            <none>
asif@master-node:~$
```


```
asif@worker02:~$ docker ps
Got permission denied while trying to connect to the Docker daemon socket at unix:///var/run/docker.sock: Get http://localhost:2375/jsoninfo: dial unix /var/run/docker.sock: connect: permission denied
asif@worker02:~$ sudo docker ps
[sudo] password for asif:
CONTAINER ID   IMAGE     PORTS   NAMES
3635422a4e88   nginx    80/tcp   k8s_master-ac-d
Up 13 minutes
da0-4722-91b3-8047656aa70b_0
```

- See what is inside a pod: kubectl describe pod master-ac-demo

What is a namespace in Kubernetes?

- A namespace is an abstracted domain where different teams can deploy their own pods
- One team cannot see other team's Pods

```
aslf@master-node:~$ kubectl get namespaces
NAME                STATUS    AGE
default             Active    13h
kube-node-lease     Active    13h
kube-public         Active    13h
kube-system         Active    13h
aslf@master-node:~$ kubectl get pods -A
NAMESPACE   NAME                                     READY   STATUS    RESTARTS   AGE
default     master-ac-demo                         1/1     Running   0           71s
kube-system coredns-558bd4d5db-8qx4b              1/1     Running   0           13h
kube-system coredns-558bd4d5db-q25jn              1/1     Running   0           13h
kube-system etcd-master-node                  1/1     Running   1           13h
kube-system kube-apiserver-master-node  1/1     Running   1           13h
kube-system kube-controller-manager-master-node  1/1     Running   1           13h
kube-system kube-flannel-ds-jkx17       1/1     Running   1           13h
kube-system kube-flannel-ds-k2ff5       1/1     Running   1           13h
kube-system kube-flannel-ds-vsmc6       1/1     Running   1           13h
kube-system kube-proxy-bh4tg            1/1     Running   1           13h
kube-system kube-proxy-mslft            1/1     Running   1           13h
kube-system kube-proxy-nhj65            1/1     Running   1           13h
kube-system kube-scheduler-master-node  1/1     Running   1           13h
kube-system weave-net-jz7vt             2/2     Running   3           13h
kube-system weave-net-s2r8k             2/2     Running   3           13h
kube-system weave-net-v9v5z             2/2     Running   3           13h
aslf@master-node:~$
```



All Cluster components are
inside kube-system namespace

Creating Pod from Configuration - Declarative

- Creating Pod using YAML file is called declarative way of defining resources
- In production it is NOT advised to run imperative commands to deploy a pod using the “kubectl run” command. Instead we use YAML or JSON file
- Complete Pod specifications should be written in a YAML file
 - Vim podSpec.yaml

We can define:

1. API version
2. Resource
3. Information
4. Specification

```
apiVersion: v1
kind: Pod
metadata:
  name: mypod
spec:
  containers:
  - image: nginx
    name: mynginx
```

Define different names
for Pod and container
inside the Pod

Execute YAML file:
kubectl create -f podSpec.yaml
kubectl describe pod mypod

Creating Pod from Template

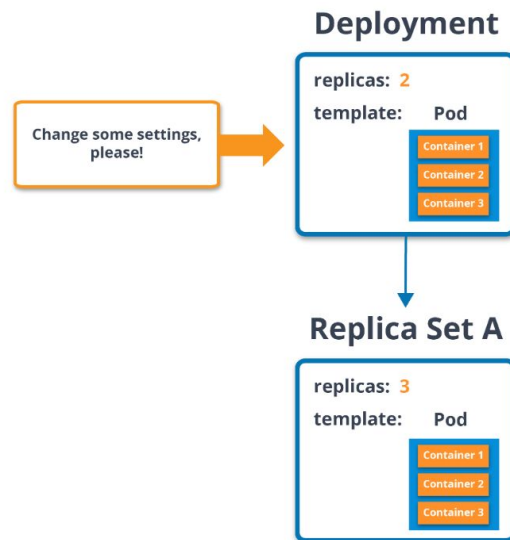
- Create a Template for Pods from imperative commands with dry run
 - `kubectl run newpod --image=nginx --dry-run -o yaml`

```
aslf@master-node:~$ kubectl run newpod --image=nginx --dry-run -o yaml
W0704 16:57:34.667715 2541 helpers.go:557] --dry-run is deprecated and can be replaced with --dry-run=client.
apiVersion: v1
kind: Pod
metadata:
  creationTimestamp: null
  labels:
    run: newpod
  name: newpod
spec:
  containers:
  - image: nginx
    name: newpod
    resources: {}
  dnsPolicy: ClusterFirst
  restartPolicy: Always
status: {}
```

- Output the printout in a file
 - `kubectl run newpod --image=nginx --dry-run -o yaml > newpodspec.yaml`

Deployment in Kubernetes

- There is a problem with just running the yaml file. There is a possibility that someone can delete that pod and it will not run again. High availability is not yet established. We use deployment to solve this problem
- Deployment can be defined to create new replica sets
- It can be defined to remove the existing deployment and use all their resources with new deployments
- Selector field defines how the pods management sequence is determined by deployment



Deployment in Kubernetes (cont)

- Create a deployment using the following command
 - `kubectl create deployment nginx --image=nginx --dry-run -o yaml > deployment.yaml`
 - `kubectl create -f deployment.yaml`
 - `kubectl get deployments.apps nginx -o wide`
 - `kubectl get pods`

```
asif@master-node:~$ kubectl get deployments.apps nginx -o wide
NAME    READY   UP-TO-DATE   AVAILABLE   AGE    CONTAINERS   IMAGES   SELECTOR
nginx   2/2     2            2           3m50s   nginx        nginx    app=nginx
asif@master-node:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
master-ac-demo                      1/1     Running   0           111m
mypod                               1/1     Running   0           64m
nginx-7848d4b86f-9vb5h              1/1     Running   0           4m6s
nginx-7848d4b86f-xtm9p              1/1     Running   0           4m6s
```

```
apiVersion: apps/v1
kind: Deployment
metadata:
  creationTimestamp: null
  labels:
    app: nginx
  name: nginx
spec:
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: nginx
    spec:
      containers:
      - image: nginx
        name: nginx
        ports:
        - containerPort: 80
```

Service in Kubernetes

- Now that we have created a deployment or launched an application inside our cluster, how do we access the application?
 - `kubectl expose deployment nginx --port=80 --target-port=80 --type=NodePort`
 - OR `kubectl create service NodePort --tcp=80:80`
 - `kubectl get service`
 - `kubectl get pods -o wide`

```
asif@master-node:~$ kubectl expose deployment nginx --port=80 --target-port=80 --type=NodePort
service/nginx exposed
asif@master-node:~$ kubectl get service
NAME         TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)          AGE
kubernetes   ClusterIP   10.96.0.1     <none>       443/TCP          16h
nginx        NodePort    10.97.185.213 <none>       80:32136/TCP     51s
asif@master-node:~$ kubectl get pods -o wide
NAME                READY   STATUS    RESTARTS   AGE   IP            NODE       NOMINATED NODE   READINESS GATES
master-ac-demo      1/1     Running   0          128m  10.244.2.3    worker02   <none>           <none>
mypod               1/1     Running   0          81m   10.244.2.4    worker02   <none>           <none>
nginx-7848d4b86f-9vb5h  1/1     Running   0          20m   10.244.2.5    worker02   <none>           <none>
nginx-7848d4b86f-xtm9p  1/1     Running   0          20m   10.244.1.3    worker01   <none>           <none>
```

- Access the application from a browser, type: ip:port “ex: 10.50.1.30:32136”

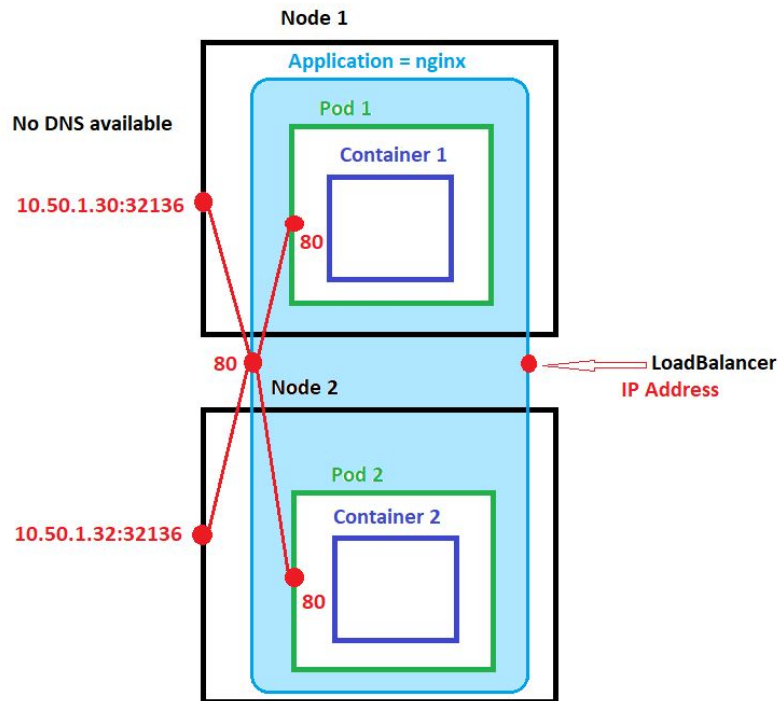


Service in Kubernetes (cont)

- NodePort is not the ideal way for launching an application on the Kubernetes cluster.
- The user needs to know the IP address of each node inside the cluster which is not practical
- The solution is to use LoadBalancer instead of NodePort
- However Nodeport solves the problem of availability. How can we test it
 - `kubectl delete pod nginx-7848d4b86f-9vb5h`
 - `kubectl get pods`

```
asif@master-node:~$ kubectl delete pod nginx-7848d4b86f-9vb5h
pod "nginx-7848d4b86f-9vb5h" deleted
asif@master-node:~$ kubectl get pods
```

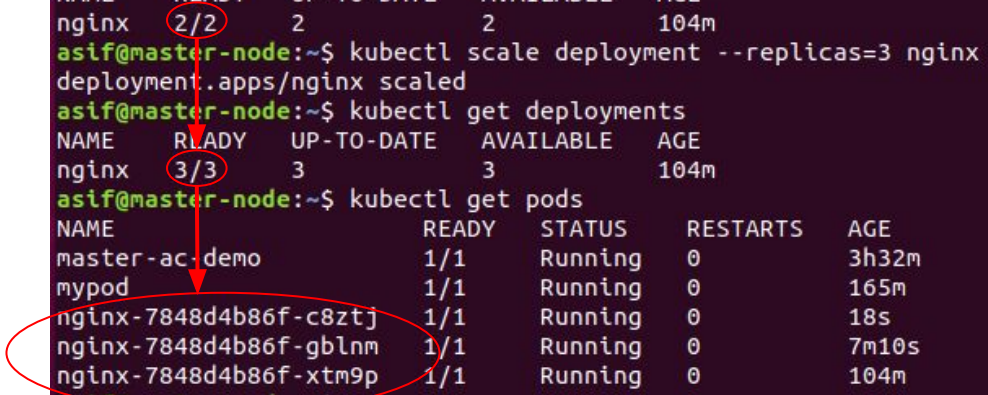
NAME	READY	STATUS	RESTARTS	AGE
master-ac-demo	1/1	Running	0	166m
mypod	1/1	Running	0	119m
nginx-7848d4b86f-ndznv	1/1	Running	0	15s
nginx-7848d4b86f-xtm9p	1/1	Running	0	58m



Scaling Using Replica Sets

- In order to scale up instances/pods we use the following command:
 - `kubectl scale deployment --replicas=3 nginx`
 - `kubectl get deployments`
 - `kubectl get pods`

```
asif@master-node:~$ kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginx     2/2     2            2           104m
asif@master-node:~$ kubectl scale deployment --replicas=3 nginx
deployment.apps/nginx scaled
asif@master-node:~$ kubectl get deployments
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
nginx     3/3     3            3           104m
asif@master-node:~$ kubectl get pods
NAME                                READY   STATUS    RESTARTS   AGE
master-ac-demo                      1/1     Running   0           3h32m
mypod                               1/1     Running   0           165m
nginx-7848d4b86f-c8ztj              1/1     Running   0           18s
nginx-7848d4b86f-gblnm              1/1     Running   0           7m10s
nginx-7848d4b86f-xtm9p              1/1     Running   0           104m
```





Kubernetes Rolling Update

- Suppose we have a new version of our application “nginx”
- How do we update the application without any service disruption?
- Kubernetes does this automatically by using the following command:
 - `kubectl set image deploy/nginx nginx=nginx:1.9.1`

```
asif@master-node:~$ kubectl set image deploy/nginx nginx=nginx:1.9.1
deployment.apps/nginx image updated
asif@master-node:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
master-ac-demo	1/1	Running	0	3h59m
mypod	1/1	Running	0	3h11m
nginx-694854bbcb-462l6	0/1	ContainerCreating	0	4s
nginx-7848d4b86f-c8ztj	1/1	Running	0	27m
nginx-7848d4b86f-gblnm	1/1	Running	0	33m
nginx-7848d4b86f-xtm9p	1/1	Running	0	131m

```
asif@master-node:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
master-ac-demo	1/1	Running	0	3h59m
mypod	1/1	Running	0	3h12m
nginx-694854bbcb-462l6	0/1	ContainerCreating	0	11s
nginx-7848d4b86f-c8ztj	1/1	Running	0	27m
nginx-7848d4b86f-gblnm	1/1	Running	0	34m
nginx-7848d4b86f-xtm9p	1/1	Running	0	131m

```
asif@master-node:~$ kubectl get pods
```

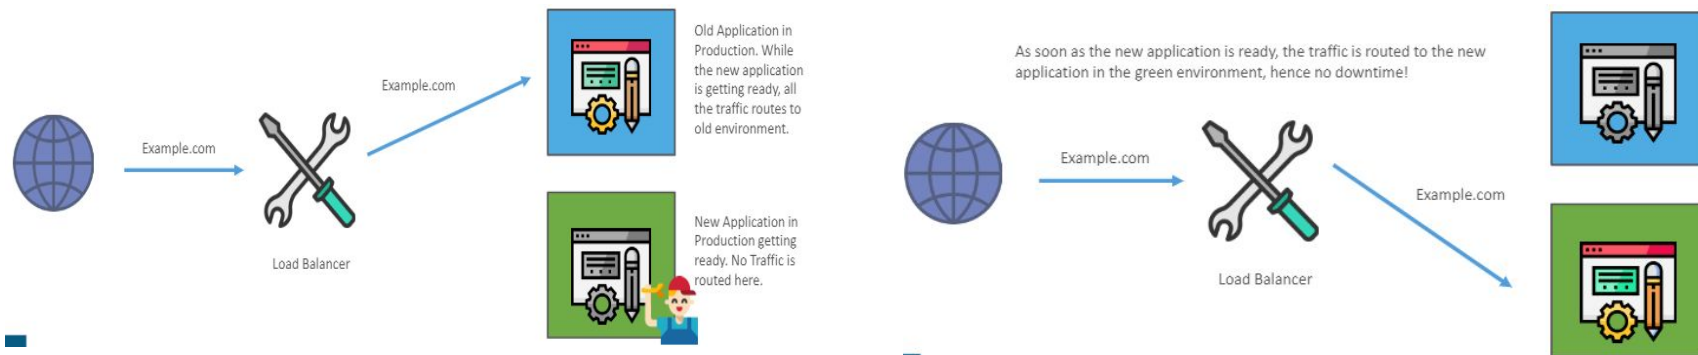
NAME	READY	STATUS	RESTARTS	AGE
master-ac-demo	1/1	Running	0	3h59m
mypod	1/1	Running	0	3h12m
nginx-694854bbcb-462l6	1/1	Running	0	21s
nginx-694854bbcb-d4wv9	0/1	ContainerCreating	0	5s
nginx-7848d4b86f-gblnm	1/1	Running	0	34m
nginx-7848d4b86f-xtm9p	1/1	Running	0	131m

```
asif@master-node:~$ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
master-ac-demo	1/1	Running	0	3h59m
mypod	1/1	Running	0	3h12m
nginx-694854bbcb-462l6	1/1	Running	0	49s
nginx-694854bbcb-d4wv9	1/1	Running	0	33s
nginx-694854bbcb-f7tgm	1/1	Running	0	17s

Blue Green Deployment Model

- Blue-green deployment is a technique that reduces downtime and risk by running two identical production environments called Blue and Green
- The updated application gets setup in the new environment (Green), while old application remains in its own environment (Blue) untouched
- Traffic stays with the blue environment until the green environment is ready
- As soon as the new application is ready, the traffic is routed to the new application in the green environment, therefore there is no downtime



End of Part 2

